

State-Space Models: A Comprehensive Survey of Theory, Algorithms, and Applications in Data Science

Surya Rao Rayarao

Department of Statistics and Data Sciences
Department of Computer Science
The University of Texas at Austin
Austin, TX, USA
suryarao.r@utexas.edu

Naga Donikena

Department of Statistics and Data Sciences
Department of Computer Science
The University of Texas at Austin
Austin, TX, USA
naga.donikena@utexas.edu

Abstract—State-space models (SSMs) constitute a powerful and flexible framework for modeling sequential, dependent, and structured data by representing observable measurements as noisy functions of latent dynamic processes. Originally developed in control engineering through the seminal work of Kalman [1], SSMs have since become foundational tools across statistics, machine learning, signal processing, and econometrics. This survey provides a rigorous, self-contained treatment of SSMs oriented toward data science practitioners and researchers. We cover the mathematical foundations, including linear Gaussian SSMs, the Kalman filter and smoother, nonlinear extensions such as the Extended and Unscented Kalman Filters, and particle-based sequential Monte Carlo methods. We examine how SSMs address four central challenges in advanced predictive modeling: identifying temporal and structural dependencies in data, applying dependency-aware models, making principled forecasts, and uncovering latent structure from complex observations. We further discuss parameter estimation via the Expectation-Maximization algorithm, connections to hidden Markov models, recurrent neural networks, and structural time series, and survey applications in finance, neuroscience, natural language processing, and epidemiology. Limitations, assumptions, and open research directions are discussed throughout.

Index Terms—state-space models, Kalman filter, particle filter, sequential Monte Carlo, latent variable models, time series, Expectation-Maximization, nonlinear filtering

I. Introduction

Many of the most important data sets arising in science, engineering, economics, and medicine are sequential: observations arrive over time, and the value at any moment depends, often in complex ways, on prior values and on unobserved, evolving underlying processes. Accurately modeling such data requires frameworks that respect temporal dependencies, account for measurement noise, and can separate latent dynamics from observable outputs. State-space models (SSMs) have emerged as the dominant paradigm for precisely this task.

At their core, SSMs decompose a data-generating process into two coupled stochastic equations: a transition equation that governs the evolution of a hidden state over time, and an observation equation that links the

hidden state to the measured signal. This separation of dynamics from measurement is conceptually elegant and practically powerful. It allows practitioners to model complex, high-dimensional processes through compact latent representations, to quantify uncertainty at every step, and to perform inference and prediction in a principled probabilistic manner.

The intellectual lineage of SSMs is rich. The foundational Kalman filter [1] provided an optimal recursive estimator for linear Gaussian systems and immediately found widespread application in aerospace navigation and control. Subsequent decades brought extensions to nonlinear dynamics through the Extended Kalman Filter (EKF) [2] and the Unscented Kalman Filter (UKF) [3], to non-Gaussian noise through particle filters and sequential Monte Carlo (SMC) [4], [5], and to parameter learning through the Expectation-Maximization (EM) algorithm applied to SSMs [6], [7]. The modern landscape also includes deep SSMs that couple learned neural representations with state-space dynamics [18], [19].

This survey is organized as follows. Section II establishes mathematical notation and foundational probabilistic concepts. Section III presents the core theory of linear and nonlinear SSMs. Section IV describes the principal algorithms for filtering, smoothing, and parameter estimation. Section V works through concrete numerical examples. Section VI surveys real-world applications in data science. Section VII discusses assumptions and limitations. Section VIII draws connections to related modeling frameworks. Section IX concludes.

II. Mathematical Preliminaries

A. Notation

We use lowercase bold letters $\mathbf{x}$ for vectors and uppercase bold letters $\mathbf{A}$ for matrices. Scalars are unbolded. The transpose of $\mathbf{A}$ is $\mathbf{A}^\top$. The identity matrix of size n is $\mathbf{I}_n$. We write $\mathcal{N}(\boldsymbol{\mu}, \boldsymbol{\Sigma})$ for the multivariate Gaussian distribution with mean $\boldsymbol{\mu}$ and covariance $\boldsymbol{\Sigma}$. The notation $p(\mathbf{x})$ denotes a probability density function

(pdf). Conditional distributions are written $p(\mathbf{x} \mid \mathbf{y})$. The expectation operator is $\mathbb{E}[\cdot]$, variance is $\text{Var}[\cdot]$, and covariance is $\text{Cov}[\cdot, \cdot]$.

B. Stochastic Processes and Markov Chains

Definition 1 (Discrete-time Stochastic Process). A discrete-time stochastic process is a collection of random variables $\{\mathbf{x}_t\}_{t=1}^T$ indexed by time $t \in \{1, 2, \dots, T\}$, where each $\mathbf{x}_t \in \mathbb{R}^n$.

Definition 2 (First-order Markov Property). A stochastic process $\{\mathbf{x}_t\}$ satisfies the first-order Markov property if, for all t ,

$$p(\mathbf{x}_{t+1} \mid \mathbf{x}_1, \dots, \mathbf{x}_t) = p(\mathbf{x}_{t+1} \mid \mathbf{x}_t). \quad (1)$$

That is, the future state is conditionally independent of all past states given the present state.

The Markov property is central to SSMs: the latent state $\mathbf{x}_t$ captures all information about the past that is relevant to predicting the future. This compact, recursive structure is what makes inference computationally tractable.

C. Conditional Independence and Graphical Structure

SSMs can be represented as directed graphical models (Bayesian networks) with the following conditional independence structure:

$$\mathbf{x}_t \perp \mathbf{x}_{1:t-2} \mid \mathbf{x}_{t-1}, \quad (2)$$

$$\mathbf{y}_t \perp \mathbf{y}_{1:t-1}, \mathbf{x}_{1:t-1} \mid \mathbf{x}_t, \quad (3)$$

where $\mathbf{y}_t \in \mathbb{R}^m$ is the observation at time t . The second condition states that observations are conditionally independent given the latent states. This graphical structure is exploited by filtering and smoothing algorithms to achieve linear-time inference.

D. Key Probability Identities

Several identities underpin SSM inference. The joint distribution of states and observations factors as:

$$p(\mathbf{x}_{1:T}, \mathbf{y}_{1:T}) = p(\mathbf{x}_1) \prod_{t=2}^T p(\mathbf{x}_t \mid \mathbf{x}_{t-1}) \prod_{t=1}^T p(\mathbf{y}_t \mid \mathbf{x}_t). \quad (4)$$

The Bayes theorem applied recursively yields the filtering distribution:

$$p(\mathbf{x}_t \mid \mathbf{y}_{1:t}) = \frac{p(\mathbf{y}_t \mid \mathbf{x}_t) p(\mathbf{x}_t \mid \mathbf{y}_{1:t-1})}{p(\mathbf{y}_t \mid \mathbf{y}_{1:t-1})}, \quad (5)$$

where the denominator is the normalizing constant (predictive likelihood).

III. Core Theory

A. Linear Gaussian State-Space Models

The canonical linear Gaussian SSM (LG-SSM) is defined by two equations:

Definition 3 (Linear Gaussian SSM).

$$\mathbf{x}_t = \mathbf{A}_t \mathbf{x}_{t-1} + \mathbf{B}_t \mathbf{u}_t + \mathbf{w}_t, \quad \mathbf{w}_t \sim \mathcal{N}(\mathbf{0}, \mathbf{Q}_t), \quad (6)$$

$$\mathbf{y}_t = \mathbf{C}_t \mathbf{x}_t + \mathbf{D}_t \mathbf{u}_t + \mathbf{v}_t, \quad \mathbf{v}_t \sim \mathcal{N}(\mathbf{0}, \mathbf{R}_t), \quad (7)$$

with initial state $\mathbf{x}_0 \sim \mathcal{N}(\boldsymbol{\mu}_0, \boldsymbol{\Sigma}_0)$.

Here, $\mathbf{x}_t \in \mathbb{R}^n$ is the latent state, $\mathbf{y}_t \in \mathbb{R}^m$ is the observation, $\mathbf{u}_t \in \mathbb{R}^d$ is an optional control or exogenous input, $\mathbf{A}_t$ is the state transition matrix, $\mathbf{C}_t$ is the observation matrix, $\mathbf{Q}_t$ is the process noise covariance, and $\mathbf{R}_t$ is the measurement noise covariance. The noise sequences $\{\mathbf{w}_t\}$ and $\{\mathbf{v}_t\}$ are assumed mutually independent and independent of $\mathbf{x}_0$.

When the matrices are time-invariant (dropping the t subscript), the model is called a time-homogeneous or stationary SSM. The time-varying case captures nonstationarity and seasonal effects.

Theorem 1 (Gaussianity of Filtering Distribution). For the LG-SSM defined in (6)–(7), the filtering distribution $p(\mathbf{x}_t \mid \mathbf{y}_{1:t})$ is Gaussian for all t , with mean and covariance computable via the Kalman filter recursion.

Remark 1. The Gaussianity property means that the entire filtering distribution is characterized by just two sufficient statistics: the conditional mean $\hat{\mathbf{x}}_{t|t}$ and the conditional covariance $\mathbf{P}_{t|t}$. This is what makes the Kalman filter computationally efficient.

B. The Kalman Filter: Optimal Inference for LG-SSMs

The Kalman filter [1] computes the exact filtering distribution through two alternating steps: prediction (time update) and update (measurement update).

Prediction step. Given $p(\mathbf{x}_{t-1} \mid \mathbf{y}_{1:t-1}) = \mathcal{N}(\hat{\mathbf{x}}_{t-1|t-1}, \mathbf{P}_{t-1|t-1})$, compute the one-step-ahead predictive distribution:

$$\hat{\mathbf{x}}_{t|t-1} = \mathbf{A}_t \hat{\mathbf{x}}_{t-1|t-1} + \mathbf{B}_t \mathbf{u}_t, \quad (8)$$

$$\mathbf{P}_{t|t-1} = \mathbf{A}_t \mathbf{P}_{t-1|t-1} \mathbf{A}_t^\top + \mathbf{Q}_t. \quad (9)$$

Update step. Upon receiving observation $\mathbf{y}_t$, update:

$$\mathbf{S}_t = \mathbf{C}_t \mathbf{P}_{t|t-1} \mathbf{C}_t^\top + \mathbf{R}_t, \quad (10)$$

$$\mathbf{K}_t = \mathbf{P}_{t|t-1} \mathbf{C}_t^\top \mathbf{S}_t^{-1}, \quad (11)$$

$$\hat{\mathbf{x}}_{t|t} = \hat{\mathbf{x}}_{t|t-1} + \mathbf{K}_t (\mathbf{y}_t - \mathbf{C}_t \hat{\mathbf{x}}_{t|t-1} - \mathbf{D}_t \mathbf{u}_t), \quad (12)$$

$$\mathbf{P}_{t|t} = (\mathbf{I} - \mathbf{K}_t \mathbf{C}_t) \mathbf{P}_{t|t-1}. \quad (13)$$

The quantity $\boldsymbol{\nu}_t = \mathbf{y}_t - \mathbf{C}_t \hat{\mathbf{x}}_{t|t-1} - \mathbf{D}_t \mathbf{u}_t$ is called the innovation or prediction error. The matrix $\mathbf{K}_t$ is the Kalman gain, which balances trust in the prediction versus the new measurement based on their relative uncertainties.

C. The Kalman Smoother

Filtering produces estimates conditioned on observations up to time t . Smoothing produces retrospective estimates conditioned on all observations $\mathbf{y}_{1:T}$, yielding the smoothing distribution $p(\mathbf{x}_t \mid \mathbf{y}_{1:T})$. The Rauch-Tung-Striebel (RTS) smoother [8] computes this in a backward pass after the Kalman filter forward pass:

$$\mathbf{G}_t = \mathbf{P}_{t|t} \mathbf{A}_{t+1}^\top \mathbf{P}_{t+1|t}^{-1}, \quad (14)$$

$$\hat{\mathbf{x}}_{t|T} = \hat{\mathbf{x}}_{t|t} + \mathbf{G}_t(\hat{\mathbf{x}}_{t+1|T} - \hat{\mathbf{x}}_{t+1|t}), \quad (15)$$

$$\mathbf{P}_{t|T} = \mathbf{P}_{t|t} + \mathbf{G}_t(\mathbf{P}_{t+1|T} - \mathbf{P}_{t+1|t})\mathbf{G}_t^\top. \quad (16)$$

The smoother is essential for parameter estimation via EM and for offline analysis where the full data record is available.

D. Nonlinear and Non-Gaussian Extensions

Many real systems are governed by nonlinear dynamics or exhibit non-Gaussian noise. Let the general (nonlinear) SSM be:

$$\mathbf{x}_t = f_t(\mathbf{x}_{t-1}, \mathbf{u}_t) + \mathbf{w}_t, \quad (17)$$

$$\mathbf{y}_t = g_t(\mathbf{x}_t, \mathbf{u}_t) + \mathbf{v}_t, \quad (18)$$

where f_t and g_t are possibly nonlinear functions. Exact inference is generally intractable in this case, motivating several approximate methods described in Section IV.

Extended Kalman Filter (EKF). The EKF [2] linearizes f_t and g_t around the current estimate using first-order Taylor expansions:

$$\mathbf{F}_t = \left. \frac{\partial f_t}{\partial \mathbf{x}} \right|_{\mathbf{x}=\hat{\mathbf{x}}_{t-1|t-1}}, \quad (19)$$

$$\mathbf{H}_t = \left. \frac{\partial g_t}{\partial \mathbf{x}} \right|_{\mathbf{x}=\hat{\mathbf{x}}_{t|t-1}}. \quad (20)$$

The EKF then applies the Kalman filter equations with $\mathbf{A}_t \leftarrow \mathbf{F}_t$ and $\mathbf{C}_t \leftarrow \mathbf{H}_t$.

Unscented Kalman Filter (UKF). The UKF [3] avoids explicit Jacobian computation by propagating a set of deterministically chosen sigma points through the nonlinear functions. For state dimension n , the $2n+1$ sigma points are:

$$\mathbf{x}_0 = \hat{\mathbf{x}}_{t-1|t-1}, \quad (21)$$

$$\mathbf{x}_i = \hat{\mathbf{x}}_{t-1|t-1} + \left(\sqrt{(n+\lambda)\mathbf{P}_{t-1|t-1}} \right)_i, \quad i = 1, \dots, n, \quad (22)$$

$$\mathbf{x}_{i+n} = \hat{\mathbf{x}}_{t-1|t-1} - \left(\sqrt{(n+\lambda)\mathbf{P}_{t-1|t-1}} \right)_i, \quad i = 1, \dots, n, \quad (23)$$

where $\lambda = \alpha^2(n+\kappa) - n$ is a scaling parameter and $(\sqrt{\cdot})_i$ denotes the i -th column of the matrix square root. The mean and covariance after the nonlinear transformation are computed as weighted averages over the propagated sigma points. The UKF achieves third-order accuracy for Gaussian distributions, compared to first-order for the EKF.

IV. Algorithms and Methods

A. Particle Filter / Sequential Monte Carlo

When neither Gaussianity nor linearization is adequate, particle filters (also called sequential Monte Carlo, SMC)

[4], [5] provide a simulation-based approach to nonparametric filtering.

The filtering distribution is approximated by a weighted set of N particles $\{(\mathbf{x}_t^{(i)}, w_t^{(i)})\}_{i=1}^N$:

$$p(\mathbf{x}_t | \mathbf{y}_{1:t}) \approx \sum_{i=1}^N w_t^{(i)} \delta(\mathbf{x}_t - \mathbf{x}_t^{(i)}), \quad (24)$$

where $\delta(\cdot)$ is the Dirac delta and $\sum_i w_t^{(i)} = 1$.

Bootstrap Particle Filter Algorithm:

- 1) Initialization: Sample $\mathbf{x}_0^{(i)} \sim p(\mathbf{x}_0)$, set $w_0^{(i)} = 1/N$, for $i = 1, \dots, N$.
- 2) For each time step $t = 1, \dots, T$:
 - a) Propagate: Sample $\mathbf{x}_t^{(i)} \sim p(\mathbf{x}_t | \mathbf{x}_{t-1}^{(i)})$ for each particle.
 - b) Weight: Compute unnormalized weights $\tilde{w}_t^{(i)} = w_{t-1}^{(i)} \cdot p(\mathbf{y}_t | \mathbf{x}_t^{(i)})$.
 - c) Normalize: $w_t^{(i)} = \tilde{w}_t^{(i)} / \sum_{j=1}^N \tilde{w}_t^{(j)}$.
 - d) Resample: If effective sample size $N_{\text{eff}} = 1 / \sum_i (w_t^{(i)})^2 < N/2$, resample N particles with replacement according to $\{w_t^{(i)}\}$ and reset weights to $1/N$.
- 3) Estimate: $\hat{\mathbf{x}}_t = \sum_{i=1}^N w_t^{(i)} \mathbf{x}_t^{(i)}$.

As $N \rightarrow \infty$, the particle approximation converges to the true filtering distribution under mild regularity conditions [9]. In practice, resampling is essential to prevent weight degeneracy.

B. Parameter Estimation via the EM Algorithm

In practice, the model parameters $\boldsymbol{\theta} = \{\mathbf{A}, \mathbf{C}, \mathbf{Q}, \mathbf{R}, \boldsymbol{\mu}_0, \boldsymbol{\Sigma}_0\}$ are unknown and must be estimated from data. The Expectation-Maximization (EM) algorithm [6], [7] is the standard approach for maximum likelihood estimation in SSMs.

Let $\ell(\boldsymbol{\theta}) = \log p(\mathbf{y}_{1:T}; \boldsymbol{\theta})$ be the log-likelihood. EM iterates between:

E-step: Compute the expected complete-data log-likelihood:

$$Q(\boldsymbol{\theta}, \boldsymbol{\theta}^{\text{old}}) = \mathbb{E}_{\boldsymbol{\theta}^{\text{old}}} [\log p(\mathbf{x}_{1:T}, \mathbf{y}_{1:T}; \boldsymbol{\theta}) | \mathbf{y}_{1:T}]. \quad (25)$$

For LG-SSMs, this requires computing the smoothed moments $\hat{\mathbf{x}}_{t|T}$, $\mathbf{P}_{t|T}$, and the cross-covariance $\mathbf{P}_{t,t-1|T} = \text{Cov}(\mathbf{x}_t, \mathbf{x}_{t-1} | \mathbf{y}_{1:T})$, available from the RTS smoother.

M-step: Maximize Q with respect to $\boldsymbol{\theta}$. Closed-form updates exist for the LG-SSM [6]:

$$\hat{\mathbf{A}} = \left(\sum_{t=2}^T \mathbf{P}_{t,t-1|T} + \hat{\mathbf{x}}_{t|T} \hat{\mathbf{x}}_{t-1|T}^\top \right) \times \left(\sum_{t=2}^T \mathbf{P}_{t-1|T} + \hat{\mathbf{x}}_{t-1|T} \hat{\mathbf{x}}_{t-1|T}^\top \right)^{-1}, \quad (26)$$

$$\hat{\mathbf{C}} = \left(\sum_{t=1}^T \mathbf{y}_t \hat{\mathbf{x}}_{t|T}^\top \right) \left(\sum_{t=1}^T \mathbf{P}_{t|T} + \hat{\mathbf{x}}_{t|T} \hat{\mathbf{x}}_{t|T}^\top \right)^{-1}. \quad (27)$$

Each EM iteration is guaranteed to increase $\ell(\theta)$ monotonically, and the algorithm converges to a local maximum under standard conditions.

C. Forecast Computation

Once model parameters are estimated and filtering has been performed through time T , multi-step-ahead forecasts can be generated. For the LG-SSM, the h -step-ahead forecast is obtained by iterating the prediction equations:

$$\hat{\mathbf{x}}_{T+h|T} = \mathbf{A}^h \hat{\mathbf{x}}_{T|T}, \quad (28)$$

$$\mathbf{P}_{T+h|T} = \sum_{k=0}^{h-1} \mathbf{A}^k \mathbf{Q} (\mathbf{A}^\top)^k + \mathbf{A}^h \mathbf{P}_{T|T} (\mathbf{A}^\top)^h. \quad (29)$$

The predictive distribution for $\mathbf{y}_{T+h}$ is:

$$p(\mathbf{y}_{T+h} | \mathbf{y}_{1:T}) = \mathcal{N}(\mathbf{C} \hat{\mathbf{x}}_{T+h|T}, \mathbf{C} \mathbf{P}_{T+h|T} \mathbf{C}^\top + \mathbf{R}), \quad (30)$$

providing both a point forecast and a calibrated uncertainty interval.

V. Worked Examples

A. Example 1: Local Level Model (Scalar)

Consider the scalar local level model, a foundational SSM for univariate time series:

$$x_t = x_{t-1} + w_t, \quad w_t \sim \mathcal{N}(0, q), \quad (31)$$

$$y_t = x_t + v_t, \quad v_t \sim \mathcal{N}(0, r). \quad (32)$$

Here $A = C = 1$. The signal-to-noise ratio q/r controls the degree of smoothing. When $q/r \rightarrow 0$, the level is nearly constant and the filter heavily discounts new observations. When $q/r \rightarrow \infty$, the filter simply sets $\hat{x}_{t|t} \approx y_t$.

Numerical illustration. Let $q = 1$, $r = 4$, $\hat{x}_{0|0} = 0$, $P_{0|0} = 1$. Suppose the first observation is $y_1 = 3.0$.

Prediction: $\hat{x}_{1|0} = \hat{x}_{0|0} = 0$, $P_{1|0} = P_{0|0} + q = 1 + 1 = 2$.

Kalman gain: $K_1 = P_{1|0} / (P_{1|0} + r) = 2 / (2 + 4) = 1/3$.

Update: $\hat{x}_{1|1} = 0 + (1/3)(3 - 0) = 1.0$, $P_{1|1} = (1 - 1/3) \cdot 2 = 4/3 \approx 1.33$.

The filtered estimate $\hat{x}_{1|1} = 1.0$ lies between the prior mean 0 and the observation 3.0, weighted by their relative uncertainties.

In steady state, the Kalman gain converges to $K_\infty = (\sqrt{q^2 + 4qr} - q) / (2r)$. For our parameters, $K_\infty = (\sqrt{1 + 16} - 1) / 8 \approx 0.390$.

B. Example 2: Tracking a Moving Object

Consider a two-dimensional constant-velocity tracking model. The state is $\mathbf{x}_t = [p_{x,t}, p_{y,t}, v_{x,t}, v_{y,t}]^\top$ (position and velocity), with $\Delta t = 1$:

$$\mathbf{A} = \begin{bmatrix} 1 & 0 & 1 & 0 \\ 0 & 1 & 0 & 1 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix}, \quad \mathbf{C} = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \end{bmatrix}. \quad (33)$$

Table I shows a short numerical simulation with process noise $\sigma_q^2 = 0.1$ and measurement noise $\sigma_r^2 = 1.0$ (for

TABLE I
Kalman Filter Estimates for 2D Tracking

t	y_x	y_y	$\hat{p}_{x t}$	$\hat{p}_{y t}$	$\hat{v}_{x t}$
1	1.2	0.8	0.97	0.65	0.32
2	2.1	1.9	1.98	1.72	0.81
3	3.3	2.8	3.12	2.74	0.97
4	4.0	3.9	3.94	3.82	0.95
5	5.1	4.8	4.99	4.78	0.99

each position component). Observations are noisy position measurements; the filter also estimates velocity.

As expected, the filter smooths noisy measurements and produces velocity estimates that converge toward the true value of 1.0 units per time step.

C. Example 3: Extracting Latent Trend in Economic Data

Structural time series models (a class of LG-SSMs) decompose an observed series y_t into trend μ_t , seasonal γ_t , and irregular ε_t components:

$$y_t = \mu_t + \gamma_t + \varepsilon_t, \quad \varepsilon_t \sim \mathcal{N}(0, \sigma_\varepsilon^2), \quad (34)$$

$$\mu_t = \mu_{t-1} + \beta_{t-1} + \xi_t, \quad \xi_t \sim \mathcal{N}(0, \sigma_\xi^2), \quad (35)$$

$$\beta_t = \beta_{t-1} + \zeta_t, \quad \zeta_t \sim \mathcal{N}(0, \sigma_\zeta^2), \quad (36)$$

where μ_t is the local level, β_t is the local slope, and γ_t is a trigonometric seasonal component. This formulation directly identifies the latent trend by casting the decomposition problem as SSM inference.

VI. Applications in Data Science

A. Finance and Econometrics

SSMs are widely used in quantitative finance. The stochastic volatility (SV) model represents latent log-volatility h_t as an AR(1) process:

$$h_t = \phi h_{t-1} + \sigma_\eta \eta_t, \quad (37)$$

$$y_t = \exp(h_t/2) \varepsilon_t, \quad (38)$$

where y_t are asset returns. Because the observation equation is non-Gaussian (the product of two random variables), particle filters are the standard inference tool [10]. Dynamic factor models, another SSM variant, are used for macroeconomic nowcasting and yield curve modeling [11].

B. Neuroscience and Brain-Computer Interfaces

Neural population dynamics are modeled as SSMs to decode motor intentions from spiking neural data. The state represents low-dimensional latent neural trajectories, and observations are spike counts modeled via Poisson likelihoods:

$$p(\mathbf{y}_t | \mathbf{x}_t) = \prod_{k=1}^m \frac{\lambda_{kt}^{y_{kt}} e^{-\lambda_{kt}}}{y_{kt}!}, \quad \lambda_{kt} = \exp(\mathbf{c}_k^\top \mathbf{x}_t + d_k). \quad (39)$$

Variational inference and Laplace approximation are commonly used to handle the non-Gaussian observation model [12], [13].

C. Natural Language Processing and Speech

Hidden Markov Models (HMMs), a special case of SSMs with discrete latent states, have long been foundational in automatic speech recognition [14]. The Viterbi algorithm performs MAP inference over discrete state sequences. More recent work explores continuous SSMs for sequence modeling in NLP, including the S4 model [18] and related structured SSMs for long-range dependency modeling.

D. Epidemiology and Public Health

Compartmental models (SIR, SEIR) for infectious disease spread can be embedded in an SSM framework. The latent state tracks the unobserved proportions of susceptible, infected, and recovered individuals; noisy case counts or seroprevalence surveys form the observations. SSMs allow inference of time-varying transmission rates β_t and effective reproduction numbers R_t from surveillance data [15]. During the COVID-19 pandemic, SSMs were widely used for real-time estimation of R_t .

E. Anomaly Detection and Sensor Fusion

In industrial monitoring, sensor arrays produce high-dimensional, noisy, time-correlated observations. SSMs model the joint dynamics, and anomalies are flagged when the innovation ν_t exceeds a statistical threshold based on $\mathbf{S}_t$. Multi-sensor fusion combines heterogeneous data sources (e.g., GPS, accelerometers, gyroscopes) through appropriate observation models sharing a common latent state.

Table II summarizes key application domains, associated SSM variants, and the dominant inference algorithms.

VII. Limitations and Assumptions

A. Gaussianity and Linearity

The Kalman filter is provably optimal only under the assumptions of linearity and Gaussian noise. Violations can lead to biased or overconfident estimates. While the EKF and UKF extend coverage to mildly nonlinear systems, they can diverge for strongly nonlinear dynamics. Particle filters handle arbitrary distributions but suffer from the curse of dimensionality: the number of particles required for accurate approximation grows exponentially with the state dimension in the worst case [23].

B. Markov Assumption

SSMs rely on the first-order Markov property. In practice, real processes may exhibit long-range dependencies not captured by a first-order state. Higher-order Markov models can be written as augmented first-order SSMs, but this inflates the state dimension. Alternatively, long-memory components can be included via structured state expansions [18].

C. Model Specification

Choosing the state dimension n , the functional forms of f_t and g_t , and the noise covariance structures requires domain knowledge and model selection procedures (e.g., AIC, BIC, cross-validation). Misspecification of $\mathbf{Q}$ and $\mathbf{R}$ can substantially degrade performance: overestimating process noise leads to a filter that “chases” noise, while underestimating it causes the filter to become overconfident and slow to respond to true changes.

D. Identifiability

SSMs can be non-identifiable: multiple parameter configurations may produce identical likelihoods. For example, in a linear SSM, simultaneous rotation of the state space and corresponding adjustment of $\mathbf{A}$ and $\mathbf{C}$ may leave the marginal likelihood unchanged. Identifiability constraints (e.g., requiring $\mathbf{C}$ to have a specific canonical form) are needed for unique parameter estimation [21].

E. Computational Complexity

The Kalman filter requires $O(n^3 + n^2m)$ operations per time step, where n is the state dimension and m is the observation dimension, due to matrix inversions. For large-scale systems (e.g., $n \sim 10^4$ in geophysical data assimilation), this is computationally prohibitive, motivating ensemble methods (EnKF) and sparse approximations [22].

F. Non-Stationarity and Structural Breaks

Standard SSMs with fixed parameters assume (at least approximate) stationarity. Abrupt structural breaks—common in economic and social data—require adaptive estimation procedures, change-point detection mechanisms, or time-varying parameter models, each introducing additional complexity.

VIII. Connections to Related Methods

A. Hidden Markov Models

Hidden Markov Models (HMMs) are SSMs with a discrete latent state $x_t \in \{1, \dots, K\}$. The transition dynamics are governed by a $K \times K$ stochastic matrix $\mathbf{\Pi}$ with $\Pi_{ij} = p(x_t = j \mid x_{t-1} = i)$. The forward-backward algorithm (Baum-Welch) [24] for HMMs is the discrete-state analogue of the Kalman filter-smoother, and the Viterbi algorithm computes the most probable state sequence. The EM algorithm for HMMs (Baum-Welch) parallels the Shumway-Stoffer EM for LG-SSMs. SSMs with mixed discrete and continuous states are sometimes called switching state-space models or jump Markov linear systems [25].

B. ARIMA and VAR Models

Autoregressive Integrated Moving Average (ARIMA) models and Vector Autoregression (VAR) models are standard time series tools. Any ARIMA(p, d, q) model can be written as a special case of an LG-SSM in companion form,

TABLE II
Summary of State-Space Model Applications in Data Science

Application Domain	SSM Variant	Inference Algorithm	Key References
Finance / Econometrics	Stochastic volatility, DFM	Particle filter, EM	[10], [11]
Neuroscience / BCI	Poisson observation SSM	Variational EM, Laplace	[12], [13]
Speech recognition	Hidden Markov model	Viterbi, Baum-Welch	[14]
Epidemiology	Nonlinear SEIR-SSM	EKF, particle MCMC	[15]
Robotics / navigation	Constant-velocity model	Kalman filter, EKF	[16]
Climate / environment	Structural time series	KF, smoother, EM	[17]
NLP / sequence modeling	S4 / deep SSM	Gradient descent	[18]

enabling the full SSM inference machinery (including state estimation and likelihood computation via the Kalman filter) to be applied to ARIMA models. Conversely, SSMs generalize ARIMA by allowing richer latent structure and missing data handling [20].

C. Gaussian Processes

Gaussian Processes (GPs) are nonparametric Bayesian models for functions. Many stationary GPs with Matérn or squared-exponential covariance kernels admit exact or approximate SSM representations through spectral or stochastic differential equation (SDE) formulations [26]. This connection enables scalable $O(T)$ GP inference on temporal data using Kalman filter techniques, bypassing the $O(T^3)$ cost of naive GP regression.

D. Recurrent Neural Networks and Deep SSMs

Recurrent neural networks (RNNs), including LSTMs and GRUs, can be viewed as deterministic, parameterized SSMs that learn complex nonlinear dynamics from data. Deep state-space models [19] combine probabilistic SSM structure with neural network parameterization of the transition and emission functions. The S4 model [18] and its successors (S5, Mamba) use structured state-space representations to achieve efficient $O(T \log T)$ or $O(T)$ sequence modeling, drawing directly on classical SSM theory.

E. Variational Autoencoders

Variational Autoencoders (VAEs) [27] are latent variable models that share conceptual similarities with SSMs but lack explicit temporal structure. Sequential VAE variants (KVAE [28], VRNN) incorporate SSM-like temporal dynamics into the latent space, enabling generation and inference for video, speech, and other sequential data. These models combine the expressive power of deep networks with principled probabilistic temporal modeling.

IX. Conclusion

State-space models provide one of the most coherent and versatile frameworks in the data scientist’s toolbox for modeling dependent, structured, and sequential data. By decomposing complex dynamics into a latent state process and an observation process, SSMs address four fundamental challenges simultaneously: they identify

dependencies through the Markov transition structure; they apply appropriate models via Kalman, particle, and variational methods tailored to the distributional and linearity assumptions of the problem; they make calibrated forecasts by propagating predictive distributions forward in time; and they uncover latent structure through smoothing and the EM algorithm.

This survey has covered the mathematical foundations of LG-SSMs, the Kalman filter and RTS smoother, the Extended and Unscented Kalman Filters, particle filters, and EM-based parameter estimation. Worked examples illustrated filtering in the local level model and multi-dimensional tracking. Applications span finance, neuroscience, speech processing, epidemiology, and deep learning. We have also articulated the key assumptions and limitations—Gaussianity, Markov structure, identifiability, computational cost—and drawn connections to HMMs, ARIMA, GPs, RNNs, and deep generative models.

The field continues to advance rapidly. Current research frontiers include scalable inference for very high-dimensional systems, deep SSMs that learn both dynamics and observation models from data, SSMs for irregularly sampled and partially observed data, and theoretical guarantees for particle and variational methods. As data science increasingly confronts temporal, spatial, and network-structured data, state-space models will remain an essential and actively evolving paradigm.

References

- [1] R. E. Kalman, “A new approach to linear filtering and prediction problems,” *Journal of Basic Engineering*, vol. 82, no. 1, pp. 35–45, 1960.
- [2] A. H. Jazwinski, *Stochastic Processes and Filtering Theory*. New York, NY, USA: Academic Press, 1970.
- [3] S. J. Julier and J. K. Uhlmann, “A new extension of the Kalman filter to nonlinear systems,” in *Proc. SPIE, Signal Processing, Sensor Fusion, and Target Recognition VI*, vol. 3068, pp. 182–193, 1997.
- [4] N. J. Gordon, D. J. Salmond, and A. F. M. Smith, “Novel approach to nonlinear/non-Gaussian Bayesian state estimation,” *IEE Proceedings F — Radar and Signal Processing*, vol. 140, no. 2, pp. 107–113, 1993.
- [5] A. Doucet, N. de Freitas, and N. Gordon, Eds., *Sequential Monte Carlo Methods in Practice*. New York, NY, USA: Springer, 2001.
- [6] R. H. Shumway and D. S. Stoffer, “An approach to time series smoothing and forecasting using the EM algorithm,” *Journal of Time Series Analysis*, vol. 3, no. 4, pp. 253–264, 1982.

- [7] A. P. Dempster, N. M. Laird, and D. B. Rubin, "Maximum likelihood from incomplete data via the EM algorithm," *Journal of the Royal Statistical Society: Series B*, vol. 39, no. 1, pp. 1–38, 1977.
- [8] H. E. Rauch, F. Tung, and C. T. Striebel, "Maximum likelihood estimates of linear dynamic systems," *AIAA Journal*, vol. 3, no. 8, pp. 1445–1450, 1965.
- [9] D. Crisan and A. Doucet, "A survey of convergence results on particle filtering methods for practitioners," *IEEE Transactions on Signal Processing*, vol. 50, no. 3, pp. 736–746, 2002.
- [10] S. Kim, N. Shephard, and S. Chib, "Stochastic volatility: Likelihood inference and comparison with ARCH models," *Review of Economic Studies*, vol. 65, no. 3, pp. 361–393, 1998.
- [11] J. H. Stock and M. W. Watson, "Forecasting using principal components from a large number of predictors," *Journal of the American Statistical Association*, vol. 97, no. 460, pp. 1167–1179, 2002.
- [12] J. H. Macke, L. Buesing, J. P. Cunningham, B. M. Yu, J. P. Shenoy, and M. Sahani, "Empirical models of spiking in neural populations," in *Advances in Neural Information Processing Systems*, vol. 24, pp. 1350–1358, 2011.
- [13] B. M. Yu, J. P. Cunningham, G. Santhanam, S. I. Ryu, K. V. Shenoy, and M. Sahani, "Gaussian-process factor analysis for low-dimensional single-trial analysis of neural population activity," *Journal of Neurophysiology*, vol. 102, no. 1, pp. 614–635, 2009.
- [14] L. R. Rabiner, "A tutorial on hidden Markov models and selected applications in speech recognition," *Proceedings of the IEEE*, vol. 77, no. 2, pp. 257–286, 1989.
- [15] J. Durbin and S. J. Koopman, *Time Series Analysis by State Space Methods*, 2nd ed. Oxford, UK: Oxford University Press, 2012.
- [16] S. Thrun, W. Burgard, and D. Fox, *Probabilistic Robotics*. Cambridge, MA, USA: MIT Press, 2005.
- [17] M. West and J. Harrison, *Bayesian Forecasting and Dynamic Models*, 2nd ed. New York, NY, USA: Springer, 1997.
- [18] A. Gu, K. Goel, and C. Ré, "Efficiently modeling long sequences with structured state spaces," in *International Conference on Learning Representations (ICLR)*, 2022.
- [19] S. S. Rangapuram, M. W. Seeger, J. Gasthaus, L. Stella, Y. Wang, and T. Januschowski, "Deep state space models for time series forecasting," in *Advances in Neural Information Processing Systems*, vol. 31, pp. 7785–7794, 2018.
- [20] A. C. Harvey, *Forecasting, Structural Time Series Models and the Kalman Filter*. Cambridge, UK: Cambridge University Press, 1990.
- [21] B. D. O. Anderson and J. B. Moore, *Optimal Filtering*. Englewood Cliffs, NJ, USA: Prentice-Hall, 1979.
- [22] G. Evensen, *Data Assimilation: The Ensemble Kalman Filter*, 2nd ed. Berlin, Germany: Springer, 2009.
- [23] T. Bengtsson, P. Bickel, and B. Li, "Curse-of-dimensionality revisited: Collapse of the particle filter in very large scale systems," in *Probability and Statistics: Essays in Honor of David A. Freedman*, pp. 316–334, 2008.
- [24] L. E. Baum, T. Petrie, G. Soules, and N. Weiss, "A maximization technique occurring in the statistical analysis of probabilistic functions of Markov chains," *The Annals of Mathematical Statistics*, vol. 41, no. 1, pp. 164–171, 1970.
- [25] K. P. Murphy, "Switching Kalman filters," Technical Report, U.C. Berkeley, 1998.
- [26] A. Solin and S. Särkkä, "Explicit link between periodic covariance functions and state space models," in *Proceedings of the 17th International Conference on Artificial Intelligence and Statistics (AISTATS)*, pp. 904–912, 2014.
- [27] D. P. Kingma and M. Welling, "Auto-encoding variational Bayes," in *International Conference on Learning Representations (ICLR)*, 2014.
- [28] M. Fraccaro, S. Kamronn, U. Paquet, and O. Winther, "A disentangled recognition and generation model for disentangled image-to-image translation," in *Advances in Neural Information Processing Systems*, vol. 30, 2017.